

22-12-2021

ER-model:-

→ entities, attributes, relationships.

→ given by P.Chan in 1970's

Entities:-

→ Having physical (or) conceptual existence Ex:- student, course

→ The characteristics of entities describe its properties.

Properties for a student :- 1) Name

2) Roll No

3) Course etc.,

→ The othername of properties is 'attributes'

→ collection of entities are called 'entity sets' (or)

'entity types'

Ex:- (Sasi, --/12/2000, 420230, BTech) } Entity set.

(Arjun, --/12/2000, 420130, BTech) }

Attributes / Properties:-

→ characteristics of attributes.

→ Types:-

→ Single valued

→ Multi-valued

→ Derived

→ Simple

→ Composite

Simple attribute:-

→ Value of this attribute is single /atomic /indivisible.

Ex:- Roll No of student, course enrolled in.

Composite attribute:-

→ There will be more than one component.

Ex:- Name → first name, middle name, last name.

Single valued attribute:-

→ It has only one value

Ex:- Aadhar Number.

Multi-valued attribute:-

→ It has multiple values.

Ex:- e-mail Id

Derived attribute:-

→ Value of this attribute is derivable from another attribute.

Ex:- Age.

Diagrammatic notation:-

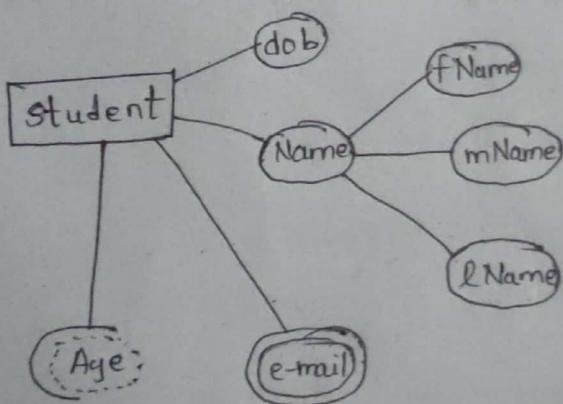
1) entity → rectangle

2) attribute → ellipse connected to rectangle

3) Composite attribute → ellipse connected to ellipse.

4) Multivalued attribute → double ellipse.

5) Derived attribute → dashed ellipse.



Domain of Attributes:-

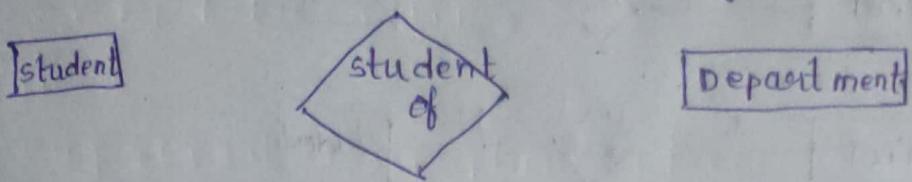
→ Domain is a set of values that an attribute can take

→ For composite attributes → domain is cross product of domain of components.

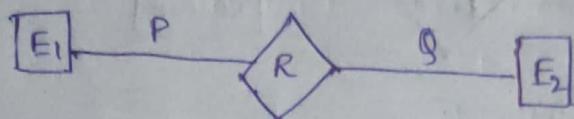
→ For multivalued attributes → set of subsets of values from basic domain

Relationship:-

- An association between entities is called relationship.
- Relationship is represented using diamond symbol.



Cardinality ratio:-



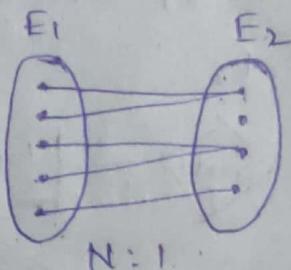
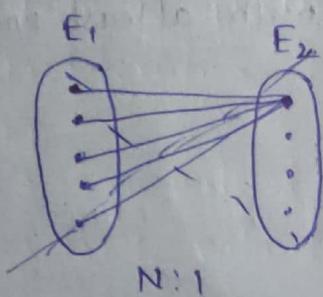
$$P = 1 \text{ or } N$$

$$Q = 1 \text{ or } N$$

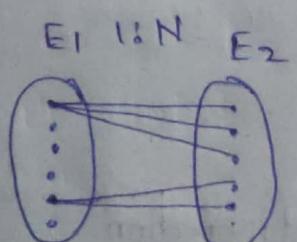
Cardinality ratio:-  $P:Q$

It can be 1) 1:1 2) 1:M 3) N:1 4) N:M

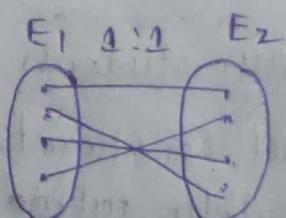
- Maximum number of entities from  $E_2$  than an entity from  $E_1$  can possibly associate through  $R$ .



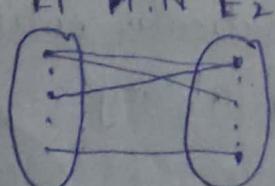
every entity in  $E_1$  with  
atmost L entity  $E_2$

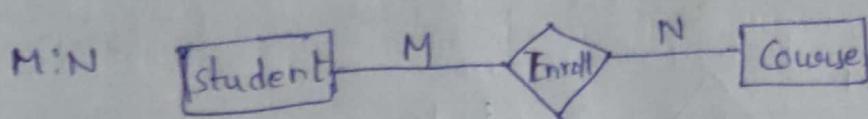
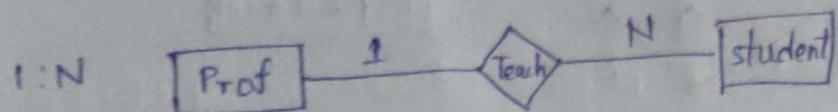
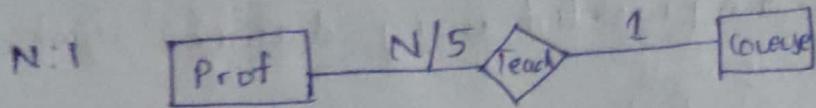
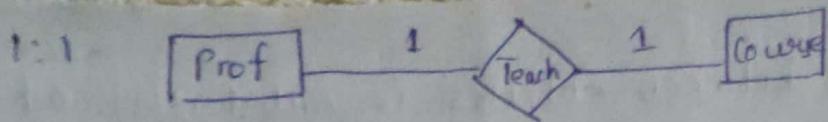


entity in  $E_1$  with  
atmost N entities in  $E_2$



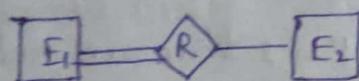
atmost 1 entity in  $E_1$   
with atmost 1 entity  
in  $E_2$ .



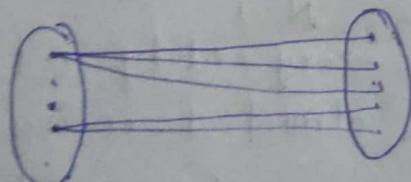


Participation Constraints:-

- Total participation (There shouldn't be any unassociated entity)
- Partial participation
- Total participation is represented by double line.



↓  
All entities in  $E_1$  should relate with atleast one entity of  $E_2$



All students are related.

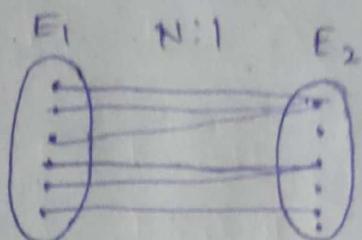
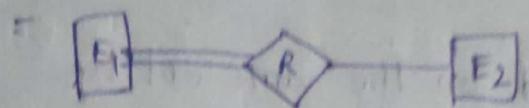
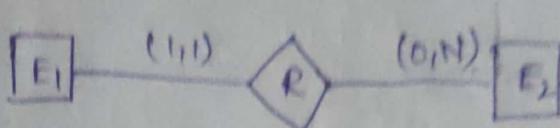
Structural Constraints:-

- Cardinality ratios along with participation constraints is structural constraints.

Min-Max Notation:-(m,n)

- m: Minimum number of times a particular entity must appear in the relationship tuples at any particular point of time.

→ n: Maximum number of times a particular entity must appear with the relationship tuples at any point of time.



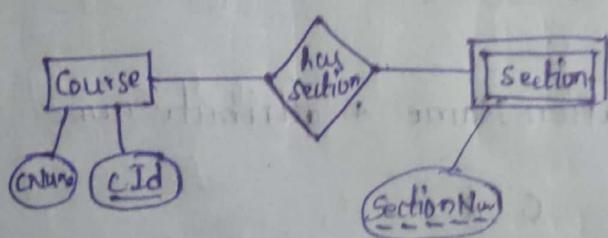
(28-12-2021)

Weak Entity sets:-

→ These sets depend for their existence on strong entity sets.

→ Weak entity is represented in double rectangle.

Here section is a weak entity.



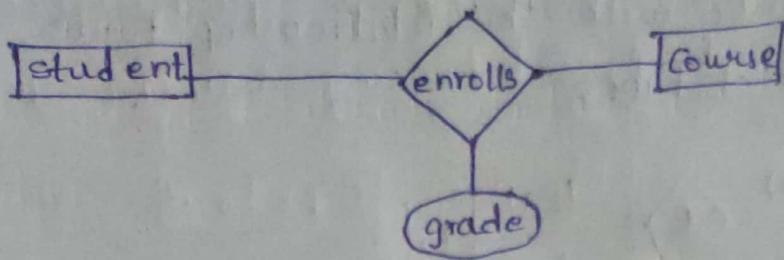
→ Weak entities have partial key i.e., as it can't serve as a complete key which helps in identifying a particular entity.

→ It should be used in combination with another key to get a particular identity.

Key - \_\_\_\_\_

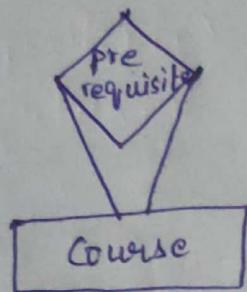
Partial key - \_\_\_\_\_

→ Relationships can also have attributes.



Recursive Relation:-

→ Relationships defined on the same entity.



Key:-

→ These are used to identify a particular entity.

ER model is a conceptual model of database.

Relational model is a logical model of database.

Relational model

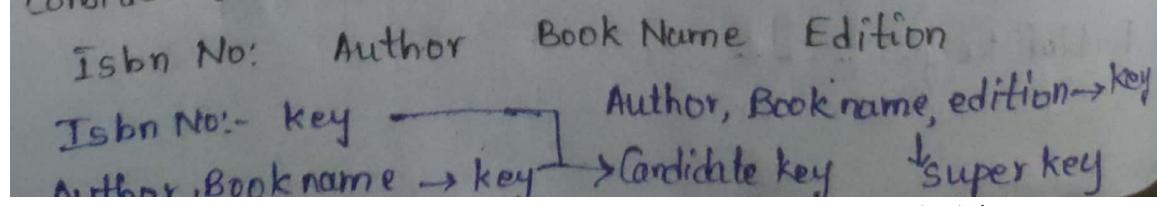
Relation schema = Relation Name + attribute name

→	S.No	S.Name	C.ID	
	1	S <sub>1</sub>	C <sub>1</sub>	
	2	S <sub>2</sub>	C <sub>2</sub>	
	3	S <sub>3</sub>	C <sub>3</sub>	

} Relation instance

→ One of the attribute (or) collection of attributes may act as a key.

→ A relation may have more than one key.  
consider the relation Book



→ One of the candidate keys is the primary key.

Minimal Super Key:-

→ The key whose subsets of it are not keys.

foreign key:-

→ A key is foreign key if it can take values of a key from another table in which it is a primary key.

Note that a table can have only one primary key.

(Referential integrity constraint)

(29-12-2021)

Relational Instance:- tuples of the relation

Database schema:-

→ collection of finite set of relations, along with integrity constraints.

Integrity constraints:-

1) Domain constraints

2) Key constraints:-

3) Referential integrity constraints:-

① values that an attribute can take is only from the declared domain.

② Its values can't be repeated for any other tuple since it should act as a key.

③ Foreign key shouldn't be violated.

Database Instance:-

→ collection of finite relational instances along with integrity constraints.

Relational Algebra:-

→ Operations to which operands are passed as relations and certain operations are performed finally another relation is obtained as output.

- set of operators
- arguments of operators are relations
- Result of Relational algebra operations is also a relation.

Operators ( $\sigma$ ,  $\Pi$ ,  $\times$ ,  $\cup$ ,  $\cap$ ,  $-$ ,  $\bowtie$ )

$\downarrow$        $\downarrow$        $\downarrow$        $\downarrow$        $\downarrow$        $\downarrow$        $\downarrow$   
 select   project   cross   union   intersect   set diff   join

$\downarrow$   
 product

Select operator ( $\sigma$ ):-

Let the relational instance of relation R be

A	B	C
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>
a <sub>3</sub>	b <sub>3</sub>	c <sub>3</sub>

$\sigma(R)$

$$A = 'a_1' \vee B = 'b_2'$$

Result:-

A	B	C
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>

$\sigma(R)$

$$A = 'a_1' \wedge B = 'b_1'$$

A      B      C

a<sub>1</sub>    b<sub>1</sub>    c<sub>1</sub>

$$\left\{ \begin{array}{l} \\ \end{array} \right. R(C,D,E) \leftarrow \sigma(R)$$

$$A = 'a_1' \wedge B = 'b_1'$$

A	D	E
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>

Project operation ( $\Pi$ ):-

$$R_1(C,D) \leftarrow \Pi_{A,B}(R)$$

Result:-

R<sub>1</sub>

C	D
a <sub>1</sub>	b <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>
a <sub>3</sub>	b <sub>3</sub>

Cross-Product operator ( $\times$ ):-

R <sub>1</sub>	a	b	c	R <sub>2</sub>	d	e	f
	a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>		d <sub>1</sub>	e <sub>1</sub>	f <sub>1</sub>
	a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>		d <sub>2</sub>	e <sub>2</sub>	f <sub>2</sub>
	a <sub>3</sub>	b <sub>3</sub>	c <sub>3</sub>		d <sub>3</sub>	e <sub>3</sub>	f <sub>3</sub>

$R_1 \times R_2$	a	b	c	d	e	f
	$a_1$	$b_1$	$c_1$	$d_1$	$e_1$	$f_1$
	$a_2$	$b_2$	$c_2$	$d_1$	$e_1$	$f_1$
	$a_3$	$b_3$	$c_3$	$d_1$	$e_1$	$f_1$
	$a_4$	$b_1$	$c_1$	$d_2$	$e_2$	$f_2$
	$a_2$	$b_2$	$c_2$	$d_2$	$e_2$	$f_2$
	$a_3$	$b_3$	$c_3$	$d_2$	$e_2$	$f_2$
	$a_4$	$b_1$	$c_1$	$d_3$	$e_3$	$f_3$
	$a_2$	$b_2$	$c_2$	$d_3$	$e_3$	$f_3$
	$a_3$	$b_3$	$c_3$	$d_3$	$e_3$	$f_3$

Union:-

→ No. of attributes should be same.

→ domains of attributes are correspondingly equal.

$R_1 \cup R_2$

a	b	c
$a_1$	$b_1$	$c_1$
$a_2$	$b_2$	$c_2$
$a_3$	$b_3$	$c_3$
$d_1$	$e_1$	$f_1$
$d_2$	$e_2$	$f_2$
$d_3$	$e_3$	$f_3$

Intersection:-

→ The tuples common to both relational instances.

$R_1$	A	B	C	$R_2$	A	B	C
	$a_1$	$b_1$	$c_1$		$a_2$	$b_2$	$c_2$
	$a_2$	$b_2$	$c_2$		$a_3$	$b_3$	$c_3$

$R_1 \cap R_2$	A	B	C
	$a_2$	$b_2$	$c_2$

Set difference:-

→ Tuples that are in first operands but not in second relational instance.

$R_1 - R_2$	A	B	C
	$a_1$	$b_1$	$c_1$

Join:  $\bowtie$

$R_1$	A	B	C
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>	
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>	

$R_2$	C	D	E
e <sub>1</sub>	d <sub>1</sub>	e <sub>1</sub>	
e <sub>2</sub>	d <sub>2</sub>	e <sub>2</sub>	

$R_1 \times R_2$

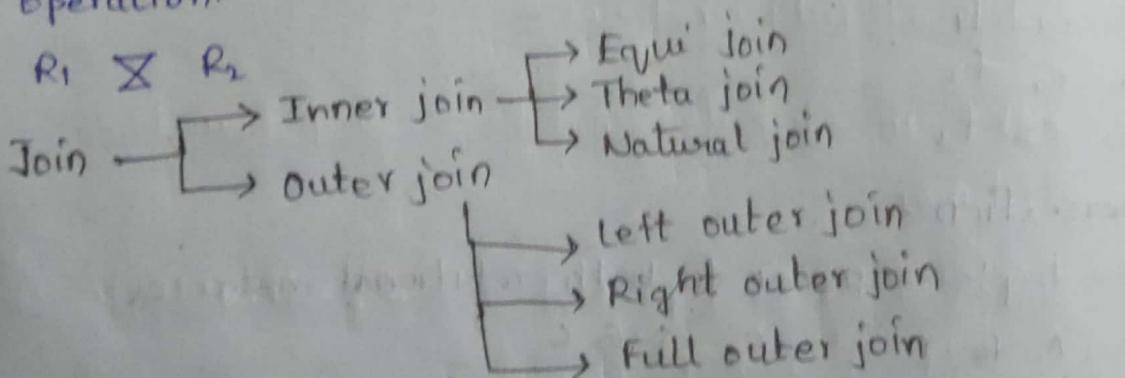
A	B	$R_1.C$	$R_2.C$	D	E
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>	e <sub>1</sub>	d <sub>1</sub>	e <sub>1</sub>
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>	e <sub>2</sub>	d <sub>2</sub>	e <sub>2</sub>
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>	e <sub>1</sub>	d <sub>1</sub>	e <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>	c <sub>2</sub>	e <sub>2</sub>	d <sub>3</sub>	e <sub>3</sub>

$\sigma (R_1 \times R_2)$

$R_1.C = R_2.C$

A	B	$R_1.C$	$R_2.C$	D	E
a <sub>1</sub>	b <sub>1</sub>	c <sub>1</sub>	c <sub>1</sub>	d <sub>1</sub>	e <sub>1</sub>

Cross-product followed by selection is called join operation.



(30-12-2021)

$R_1$	$R_2$				
A	B	C	D	E	F
a <sub>1</sub>	b <sub>1</sub>	5	8	e <sub>1</sub>	f <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>	6	9	e <sub>2</sub>	f <sub>2</sub>

$R_1$	$\bowtie$	$R_2$	CCD		
A	B	C	D	F	F
a <sub>1</sub>	b <sub>1</sub>	5	8	e <sub>1</sub>	f <sub>1</sub>
a <sub>2</sub>	b <sub>2</sub>	6	9	e <sub>2</sub>	f <sub>2</sub>

a<sub>1</sub> b<sub>1</sub> 5 9 e<sub>1</sub> f<sub>1</sub>

a<sub>2</sub> b<sub>2</sub> 6 8 e<sub>2</sub> f<sub>2</sub>

→ This join is called theta join.

→ Because conditions that are provided are called as theta.

Equi Join:-

→ In equi join, the only comparison operator should be used is 'equals to'.

Ex:- R<sub>1</sub>  $\bowtie$  R<sub>2</sub>  
c=D

Natural Join:-

→ If the two participating relations have a common attributes, then output will be those instances where attribute values are equal.

R<sub>1</sub> \* R<sub>2</sub>

R<sub>1</sub>  
A B C  
a<sub>1</sub> b<sub>1</sub> 5  
a<sub>2</sub> b<sub>2</sub> 9

R<sub>2</sub>  
D B F  
b<sub>1</sub> f<sub>1</sub>  
b<sub>2</sub> f<sub>2</sub>

R<sub>1</sub> \* R<sub>2</sub>  
A R<sub>1</sub>:B R<sub>2</sub>:C R<sub>1</sub>:C R<sub>2</sub>:B F  
a<sub>2</sub> b<sub>2</sub> 9 9 b<sub>2</sub> f<sub>2</sub>

Left outer join:-

R<sub>1</sub>  ~~$\bowtie$~~  R<sub>2</sub>

R<sub>1</sub>.C = R<sub>2</sub>.D

R<sub>1</sub>  
A B C  
a<sub>1</sub> b<sub>1</sub> 5  
a<sub>2</sub> b<sub>2</sub> 8

R<sub>2</sub>  
D E F  
5 e<sub>1</sub> f<sub>1</sub>  
8 e<sub>2</sub> f<sub>2</sub>

A B C D E F  
a<sub>1</sub> b<sub>1</sub> 5 5 e<sub>1</sub> f<sub>1</sub>  
a<sub>2</sub> b<sub>2</sub> 8 Null Null Null

→ Every tuple in the left hand side relation should be considered.

Right outer join:-

$R_1 \bowtie R_2$

$$R_1.C = R_2.D$$

$R_1$

A B C

a<sub>1</sub> b<sub>1</sub> 5

a<sub>2</sub> b<sub>2</sub> 9

$R_2$

D E F

5 e<sub>1</sub> f<sub>1</sub>

8 e<sub>2</sub> f<sub>2</sub>

A B C D E F

a<sub>1</sub> b<sub>1</sub> 5 5 e<sub>1</sub> f<sub>1</sub>

- - - 8 e<sub>2</sub> f<sub>2</sub>

Full outer join:-

$R_1 \bowtie R_2$

$$R_1.C = R_2.D$$

$R_1$

A B C

a<sub>1</sub> b<sub>1</sub> 5

a<sub>2</sub> b<sub>2</sub> 9

$R_2$

D E F

5 e<sub>1</sub> f<sub>1</sub>

8 e<sub>2</sub> f<sub>2</sub>

$R_1 \bowtie R_2$

A B C D E F

a<sub>1</sub> b<sub>1</sub> 5 5 e<sub>1</sub> f<sub>1</sub>

a<sub>2</sub> b<sub>2</sub> 9 - - -

- - -

Renaming:-

professor

Name | eid

My SQL

Oracle

Rename 'name' as 'pname'.

$\rho_{pname}(\pi_{name}(\text{Professor}))$

04-01-2022

Division:-

$S(A, B); R(A, B, C, D) \quad x \leftarrow s \times r \rightarrow X(C, D)$

S  
A      B  
a<sub>1</sub>    b<sub>1</sub>  
a<sub>2</sub>    b<sub>2</sub>

R      A      B      C      D  
a<sub>1</sub>    b<sub>1</sub>    c<sub>1</sub>    d<sub>1</sub>  
~~a<sub>2</sub>    b<sub>2</sub>    c<sub>2</sub>    d<sub>2</sub>~~  
a<sub>1</sub>    b<sub>1</sub>    c<sub>2</sub>    d<sub>2</sub>  
a<sub>2</sub>    b<sub>2</sub>    c<sub>1</sub>    d<sub>1</sub>  
a<sub>2</sub>    b<sub>2</sub>    c<sub>3</sub>    d<sub>3</sub>  
a<sub>1</sub>    b<sub>1</sub>    c<sub>2</sub>    d<sub>2</sub>

$\textcircled{1} \pi_{C, D}^r$

c <sub>1</sub>	d <sub>1</sub>	c <sub>2</sub> , d <sub>2</sub> is not in combination
c <sub>2</sub>	d <sub>2</sub>	with all tuples of S in R.
c <sub>3</sub>	d <sub>3</sub>	

$x \leftarrow r \div s$

C	D
c <sub>1</sub>	d <sub>1</sub>
c <sub>3</sub>	d <sub>3</sub>

Student [rollno] name

Enrollment [rollno] courseid

Find those students who enrolled in all of the courses offered by CSE.

1)  $\Pi_{\text{rollno}, \text{courseid}}$  (student \* enrollment)  $\rightarrow$  studEnrol

courses [courseid dno]

Department [dno dname]

2) csCourse  $\leftarrow \sigma_{\text{dname} = \text{CSE}}$  (course  $\bowtie$  department  
course dno = dept dno)

3) studEnrol  $\div$  csCourse

Aggregate functions:-

→ Avg

→ SUM

→ MIN

→ MAX

→ COUNT

→ These functions are denoted by  $\{$  (which is subscript +)

EMPLOYEE [eno] [cname] [salary] [dno]

$\{ \text{avg}(\text{salary})$  EMPLOYEE  $\rightarrow$  avg salary of employees.

grouping attributes  $\{$  aggregate functions.

$\{ \text{dno} \text{ avg}(\text{salary})$  (EMPLOYEE)

dno | avg(salary)

EMPLOYEE		dno	sex	avg(salary)
dno	sex			
avg(salary)				

- 1) Convert specification into ER diagram
- 2) Convert ER diagram into specification.

05-01-2022

S.No	S-Name	d-No	Student Mark	SQL
s <sub>1</sub>	aaa	d <sub>1</sub>	50	
s <sub>2</sub>	bbb	d <sub>1</sub>	20	
s <sub>3</sub>	ccc	d <sub>2</sub>	20	

→ attributes/columns.  
 select \* from Student; → prints whole table  
 ↓  
 Keywords

select s.No from Student; → selects only s.No column.

select sNo as id from Student; → aliases

→ The SQL language allows duplicates

If the output shouldn't have duplicates, then use distinct keyword.

select distinct d.No as dept from student;

dept  
d<sub>1</sub>  
d<sub>2</sub>

select count(\*) from Student; → gives the count of no. of rows in the table.

select sum(Marks) from Student; → gives the sum of Marks.

select \* from Student where Mark > 30;

select \* from Student where Mark > 30 and d.No = d<sub>2</sub>;

select \* from Student where Mark > 10 and Mark < 30

select \* from student where sName like '%.va%'

S.Name  
containing  
va

select \* from student where S.Name like 'va%'

start with va

select \* from student where S.Name like '%.va'

If there are two tables, student, Dept:-

the query, select \* from student, dept gives  
the cross product of two tables.

select \* from student, Dept where dNO = Dept.dnum

s.No	s.Name	d.NO	Mark	dnum	dname	phone
s <sub>1</sub>	siva	d <sub>1</sub>	10	d <sub>1</sub>	dep 1	222
s <sub>2</sub>	sivakumar	d <sub>1</sub>	20	d <sub>1</sub>	dep 1	222
s <sub>3</sub>	Vanitha	d <sub>2</sub>	30	d <sub>2</sub>	dep 2	333
s <sub>4</sub>	Rajesh	d <sub>3</sub>	50	d <sub>3</sub>	dep 3	444

select \* from student,Dept where Mark = 10

s.No	s.Name	d.No	Mark	d.Numb	d.Name	phone
s <sub>1</sub>	Siva	d <sub>1</sub>	10	d <sub>1</sub>	dep <sub>1</sub>	222
s <sub>1</sub>	siva	d <sub>1</sub>	10	d <sub>2</sub>	dep <sub>2</sub>	333
s <sub>1</sub>	siva	d <sub>1</sub>	10	d <sub>3</sub>	dep <sub>3</sub>	444

select Student \* from Student,Dept where Mark=10  
only student attributes.

Student [ S-No | S-Name | Adv ] Enrollment [ S-No | C-No | Grade ]

select S-Name from Student as S  
who S-No in (select S-No from Enrollment as E  
where student S-No = Enrollment.S-No )

Professor [ EmpId | Sex | Start Year ]

select S-No, S-Name from Student as S  
where S-Adv in (select EmpId from Professor where  
Sex = 'F')

In :- connects inner query to outer query

select S-No, S-Name from Student as S  
where S-Adv not in (select EmpId from Professor where  
Sex = 'F')

Get the (year) emp id , name of the senior most  
professor.

select EmpId, from Professor where  
Start Year ≤ All( select Start Year from Professor )

Student

S-No	S-Name	Adv	D-No
S <sub>1</sub>	aaa	e <sub>1</sub>	d <sub>1</sub> -
S <sub>2</sub>	bbb	e <sub>2</sub>	d <sub>2</sub> -
S <sub>3</sub>	ccc	e <sub>3</sub>	d <sub>3</sub> x

Professor

EmpId	Sex	Start Year
e <sub>1</sub>	M	1982
e <sub>2</sub>	M	1983

Select S-No from  
student where  
exists (select \* from  
Professor where  
EmpId = Adv )

## Student

S.No	S.Name	Adv	dNo
s <sub>1</sub>	aaa	e <sub>1</sub>	d <sub>1</sub>
s <sub>2</sub>	bbb	e <sub>2</sub>	d <sub>2</sub>
s <sub>3</sub>	ccc	e <sub>3</sub>	d <sub>1</sub>
s <sub>4</sub>	ddd	e <sub>1</sub>	d <sub>2</sub>

select S.No from student where S.adv = 'e<sub>1</sub>'

UNION

select S.No from student where S.adv = 'e<sub>2</sub>'

→ for union, attribute name and data type of attr  
should be same.

select S.No from student where dNo = 'd<sub>1</sub>'

INTERSECTION

select S.No from student where (d.No) adv = 'e<sub>1</sub>'

Output:- s<sub>1</sub>

→ Attr and data type should be same.

select S.NO from Student where 'dNo' = 'd<sub>1</sub>'

EXCEPT

select S.NO from student where adv = 'e<sub>3</sub>'

18-01-2022

Join:-

professor (pname, deptno)

Department (dname, deptId)

Select pname from Professor P where  
P.deptno = deptId from Department d and  
d.dname = "CSE"

Select \* from (Professor as P join Department  
as d on P.deptno = d.deptId) where d.dname  
= "CSE"

join = Inner join

Left outer join:-

→ All tuples from left side of join.

Select \* from (r<sub>1</sub> left outer join r<sub>2</sub> on r<sub>1</sub>.col1 =  
r<sub>2</sub>.col1) r<sub>1</sub>.col2 = 10

Right outer join:-

Select \* || from (r<sub>1</sub> right outer join r<sub>2</sub> on r<sub>1</sub>.col1 =  
r<sub>2</sub>.col1) where r<sub>1</sub>.col2 = 10

Full outer join:-

Select \* from (r<sub>1</sub> full outer join r<sub>2</sub> on r<sub>1</sub>.col1 =  
r<sub>2</sub>.col1) where r<sub>1</sub>.col2 = 10

Natural join:-

## Aggregate functions

gateMarks (reg no , name, sex , branch, city , state, marks)

Find total number of students who took gate in CS and their avg marks.

select count(\*) from gateMarks

where branch = "CS"

Select avg(marks) from gateMarks

where branch = "CS"

Select count(\*) as No:of students,

avg(marks) as Average

from gateMarks where branch = "CS"

Get max, min, Avg marks of students whose city is Hyderabad.

select max(marks) as Max,

min(marks) as Min,

avg(marks) as Avg

from gateMarks where city = "Hyderabad"

Get the names of students who obtained maximum marks in branch of EC.

select name from gateMarks where branch = "EC"

(1-01-2022)  
gateMarks (rollNo, name, sex, branch, city, state, marks)

(1, P1, m, b1, 65, 90)

(1, P2, M, b2, c1, s2, 60)

(1, P3, f, b2, c2, s2, 70)

(1, P4, m, b2, c1, s2, 60)

Select branch, max(marks) from gateMarks group by branch.

Output: (b1, 90), (b2, 70)

Select max(marks) from gateMarks group by state, sex

Select @@sql\_mode  
SET sql\_mode=select concat(@sql\_mode,  
"ONLY\_FULL\_GROUP\_BY")

i) Query: report the total enrollment in each course in the even semester of 2014, include only courses with a minimum enrollment of 10.

Enrollment (rollno, courseId, sem, year)

Select count(\*) from Enrollment

Select \* from Enrollment where sem = 2  
and year = 2014.

Select count(\*) from Enrollment where sem = 2  
and year = 2014.

Select <sup>courseId</sup> count(rollno) from Enrollment where sem = 2  
and year = 2014 group by courseId.

Select courseId, count(rollNo) from Enrollment  
where sem = 2 and year = 2014 group by  
courseId having count(rollNo) >= 10

'having' can be applied on aggregate fun  
and group by.

Views - sname, dname

select ~~sname, dname~~ from student s, department d where  
s.dno = d.dno

→ Views are virtual relations not physically  
stored, and are executed on demand.

create view stud-dpt as (select sname, dname  
from student s, department d where s.dno =  
d.dno)

→ We can treat stud-dpt as some table.

Restrictions on insertions of views:-

→ Views involving more than one table  
can't be inserted.

→ Views involving primary key of the table

→ Views involving group by clause in the  
select query.

create view ~~v1~~ v1 as (select sname from student)

delete from student where rollno = 'r2'

update student

set sname = 'lakshmi', dno = 3  
where sno = 5  
order by sname desc

25-01-2022

Functional Dependencies:-

student(studName, rollNo, sex, (deptName, hod))  
dname, hod)

Correctness of the schema:-

(s<sub>1</sub>, r<sub>1</sub>, m, d<sub>1</sub>, dn<sub>1</sub>, h<sub>1</sub>), (s<sub>2</sub>, r<sub>2</sub>, m, d<sub>1</sub>, dn<sub>2</sub>, h<sub>1</sub>)

→ For students having same dept No; dept Name and hod will be same which is repetition of data.

(null, null, null, d<sub>2</sub>, dn<sub>2</sub>, h<sub>2</sub>)

→ This is like d<sub>2</sub> has no students enrolled in it. → Insertion anomaly

→ Suppose hod of one of dept has changed. Then all values with that particular dept have to be updated with new hod which is error prone. → update anomaly, deletion anomaly

→ So, update the table with only the following.

Student(studName, rollno, sex, dno)

Department(dno, dname, hod)

Consider the relation R(x, y, z), r

X → Y : This is a functional dependency

t<sub>1</sub>, t<sub>2</sub> are any two tuples of the relation

If t<sub>1</sub>[x] = t<sub>2</sub>[x] then t<sub>1</sub>[y] = t<sub>2</sub>[y]

→ If this holds for any t<sub>1</sub>, t<sub>2</sub>, then we say

X → Y

the relation,  $R(x, y_1, y_2, y_3)$

If  $x$  is a key then  $x \rightarrow y_1, x \rightarrow y_2, x \rightarrow y_3$

Trivial functional dependencies:-

$X \rightarrow Y$

$X = \text{set of attributes}, Y = \text{set of attributes}$

If  $Y \subseteq X$ , then  $X \rightarrow Y$  is called as Trivial functional dependency.

If  $X \cap Y = \emptyset$ , then  $X \rightarrow Y$  is completely non-trivial functional dependency.

→ Using the concept of functional dependencies, Normal forms 1NF, 2NF, 3NF, BCNF, 4NF, 5NF

1-01-2022

Let  $R(x, y, z)$  be a relation

Consider a relational instance such that  $X \rightarrow Y$

$r: X \rightarrow Y$

$(x_1, y_1, z_1) t_1$

$(x_2, y_2, z_2) t_2$

$(x_1, y_1, z_3) t_3$

$\therefore X \rightarrow Y \text{ if } t_1[x] = t_2[x] \text{ then } t_1[y] = t_2[y]$

Consider  $t_1, t_3$

$t_1[x] = y_1 = t_3[x]$

Check for  $t_1[y], t_3[y]$   $t_1[y] = y_1 = t_3[y]$

$r$  satisfies the functional dependency  $X \rightarrow Y$

If  $r$  contains the following tuples:-

$t_1 (x_1, y_1, z_1)$

$t_2 (x_2, y_2, z_2)$

$t_3 (x_1, y_1, z_3)$

$t_4 (x_1, y_2, z_4)$

Here  $t_1[x] = t_3[x] = t_4[x] = x$

but  $t_1[y] = t_3[y] \neq t_4[y]$

so,  $r$  doesn't satisfy functional dependency.

If relation  $R(x, y, z, A)$  with  $x$  as a key,

then  $x \rightarrow y, x \rightarrow z, x \rightarrow A$  holds

Deriving new Functional Dependencies:-

→ Given that a set of FD's  $F$  holds on  $R$  we can infer that a certain new FD must also hold on  $R$ .

→ Ex:- given that  $x \rightarrow y, y \rightarrow z$  hold on  $R$  then  $x \rightarrow z$  must also hold.

→ Unless all FD's are known, a Relation schema is not fully specified.

Entailment Relation:-

→ F.D's  $F \vdash \{x \rightarrow y\}$  [F entails  $x \rightarrow y$ ]

if in every instance  $r$  of  $R$  on which  $F$  holds, F.D  $x \rightarrow y$  also holds.

Define  $F^+ = \{x \rightarrow y \mid F \models x \rightarrow y\}$

$F^+$  is called closure of  $F$

Armstrong's Inference Rules

1.) Reflexive rule

$$F \models \{x \rightarrow y \mid Y \subseteq X \wedge X$$

$\rightarrow$  This is called trivial FD

2.) Augmentation rule

$$\{x \rightarrow y\} \vdash \{xz \rightarrow yz\}, z \subseteq R \quad yz = x \cup z$$

3.) Transitive rule

$$\{x \rightarrow y, y \rightarrow z\} \vdash \{x \rightarrow z\}$$

4.) Decomposition or Projective rule

$$\{x \rightarrow yz\} \vdash \{x \rightarrow y\}, \{x \rightarrow yz\} \vdash \{x \rightarrow z\}$$

5.) Union (or) Additive rule

$$\{x \rightarrow y, x \rightarrow z\} \vdash \{x \rightarrow yz\}$$

6.) Pseudo transitive rule

$$\{x \rightarrow y, wy \rightarrow z\} \vdash \{wx \rightarrow z\}$$

Derive Rule 5 from 1, 2, 3

$$x \rightarrow y, x \rightarrow z$$

$$x \rightarrow y \Rightarrow x \rightarrow yx \quad \text{(Augmentation)} \\ \hookrightarrow \textcircled{1}$$

$$x \rightarrow z \Rightarrow xy \rightarrow yz \quad \text{(Augmentation)} \\ \hookrightarrow \textcircled{2}$$

$$\textcircled{1} \& \textcircled{2} \Rightarrow x \rightarrow yz \quad \text{(Transitive)}$$

Rule 4:-

$$\{x \rightarrow yz\} \vdash \{x \rightarrow y, x \rightarrow z\}$$

$$x \rightarrow yz \Rightarrow xy \rightarrow yz \Rightarrow x \rightarrow z \text{ (Augmentation)}$$

$$x \rightarrow yz \Rightarrow xz \rightarrow yz \Rightarrow x \rightarrow y \text{ (Augmentation)}$$

Rule 5:-

$$\{x \rightarrow y, wy \rightarrow z\} \vdash \{wx \rightarrow z\}$$

$$x \rightarrow y \Rightarrow wx \rightarrow wy \text{ (Augmentation)} \quad \hookrightarrow ①$$

$$wy \rightarrow z \rightarrow ②$$

$$① \& ② \Rightarrow wx \rightarrow z \text{ (transitive)}$$

→ Rules 1, 2, 3 are sufficient.

Sound and complete Inference rules:-

Sound:-

→ Every new FD  $x \rightarrow y$  derived from a given set of FD's F using Armstrong's Axioms is such that  $F \vdash \{x \rightarrow y\}$

Completeness:-

→ Any FD  $x \rightarrow y$  logically implied by F can be derived by F using Armstrong's Axioms.

Proving soundness

01-02-2022

Suppose  $x \rightarrow y$  is derived from F using AA in some n steps. If each step is correct then overall deduction would be correct.

2-02-2022

Computing  $X^+$

$R = (\text{rollNo}, \text{name}, \text{advisorId}, \text{advisorName}, \text{courseId}, \text{grade})$

$F = \{\text{rollNo} \rightarrow \text{name}, \text{rollNo} \rightarrow \text{advisorId}, \text{advisorId} \rightarrow \text{advisorName}; \text{rollNo}, \text{courseId} \rightarrow \text{grade}\}$

$\{\text{rollNo}\}^+ = \{\text{name}, \text{advisorId}, \text{advisorName}, \text{rollNo}\}$

$\{\text{advisorName}\}^+ = \{\text{advisorName}\}$

$\{\text{rollNo}, \text{courseId}\}^+ = \{\text{rollNo}, \text{name}, \text{advisorId}, \text{courseId}, \text{advisorName}, \text{grade}\}$

$\{\text{rollNo}, \text{courseId}, \text{advisorId}\}^+ = \{\text{rollNo}, \text{name}, \text{advisorId}, \text{courseId}, \text{advisorName}, \text{grade}\}$

$\rightarrow \text{rollNo}, \text{courseId}$  determines every attribute in the schema.

$\rightarrow \{\text{rollNo}, \text{courseId}\}$  is the key of R.

Normal Form

2NF:

A FD  $X \rightarrow A$  for which there is no proper subset Y of X such that  $Y \rightarrow A$  is called full functional dependency. A is said to be fully functional dependent on X.

1NF:- attributes can take only atomic values.  
(like only one value).

2NF:-

A relation schema R is in 2NF if every non-prime attribute is fully dependent.

Ex:-

Book (authorName, title, authoAffiliation, ISBN, publisher, pubYear)

Keys:- (authorName, title), ISBN

We also have authorName  $\rightarrow$  authorAffiliation

(AuthorAffiliation is not fully functional

dependent on the first key)

student (rollNo, name, dept, sex, hostelName, roomNo, admitYear)

Keys:- rollNo, (hostelName  $\rightarrow$  roomNo)

Also have hostelName  $\rightarrow$  sex

$\rightarrow$  So, this is not in 2NF

3NF:-

(07-02-2022)

$\rightarrow$  Relation schema R is in 3NF if it is in 2NF and no non prime attribute of R is transitively dependent on any key of R.

student (rollNo, name, dept,

Another definition of 3NF:-

→ Relation schema R is in 3NF if for any non-trivial FD  $X \rightarrow A$  either (i) X is superkey or (ii) A is prime.

Suppose some R violates the above definition.

BCNF:-

gradeInfo (rollNo, studName, course, grade)

Suppose:-

1) rollNo, course  $\rightarrow$  grade

2) studName, course  $\rightarrow$  grade

3) rollNo  $\rightarrow$  studName

4) studName  $\rightarrow$  rollNo

Keys:- (rollNo, course), (studName, course)

For 1,2 LHS is key. for 3,4 RHS is prime attribute.

so, gradeInfo is in 3NF.

But studName is stored redundantly along with every course being done by the student.

Boyce - Codd Normal form

→ R is in BCNF if for every non-trivial FD  $X \rightarrow A$ , X is a superkey of R.

In gradeInfo, F.D's 3, 4 are nontrivial LHS is not a superkey. So, gradeInfo is not in BCNF.

Decompose:- gradeInfo (rollNo, course, grade)  
studInfo (rollNo, studName)

→ Redundancy allowed by 3NF is disallowed by BCNF.

→ BCNF is stricter than 3NF.

→ 3NF is stricter than 2NF.

(102 - 8 sem)

## Transactions:-

→ sequence of queries performed on DB.

⇒ They a property: ACID → Durability

↓  
Isolation

↓  
Atomicity

↓  
Consistency

### Atomicity:-

→ Either all instructions in the transaction should execute or not.

→ There are no intermediate states.

### Consistency:-

→ The changes should be consistent.

### Durability:-

→ Changes made should last forever.

→ The interlinking of two or more transactions is called concurrent execution.

→ A transaction can be in any of the following states:

1) active → during execution

2) partially committed → like computations completed

3) failed → incomplete only write pending

4) Abort → terminated by the system

5) commit:

↳ finished

Active  $\rightarrow$  partially committed  $\rightarrow$  committed

$\rightarrow$  failed  $\rightarrow$  Abort

Schedule:

$\rightarrow$  A plan for concurrent execution.

let

$T_1$

read(A)

$A = A - 50$

write(A)

read(B)

$B = B + 50$

write(B)

commit

$T_2$

$T_2$

which has been lost

$T_2$

Aborted

and A = B

$T_2$

$\rightarrow$  The above is serial schedule (op's not interlinked)

$\rightarrow$  This is like plan / schedule.

read(A)

(A) then

$A = A - 50$

before

write(A)

read(A)

$temp = A * 0.1$

$A = A - temp$

write(A)

read(B)

$$B = B + 50$$

write(B)

commit

read(B)

$$B = B + \text{temp}$$

write(B)

commit

→ This is non-serial schedule

T<sub>1</sub>

read(A)

$$A = A - 50$$

T<sub>2</sub>

read(A)

$$\text{temp} = A * 0.1$$

$$A = A - \text{temp}$$

write(A)

read(B)

write(A)

read(B)

$$B = B + 50$$

write(B)

commit

$$B = B + \text{temp}$$

write(B)

commit

→ Inconsistent schedule.

08-03-2022

1)  $T_i = \text{read}(A)$   $T_j = \text{write}(A)$

2)  $T_i = \text{write}(A)$   $T_j = \text{read}(A)$

3)  ~~$T_i = \text{write}(A)$ ,  $T_j = \text{write}(B)$~~

4)  $T_i = \text{read}(A)$   $T_j = \text{read}(B) \Rightarrow$  not in conflict.

$S_1$

$T_1$

$\text{read}(A)$

$\text{write}(A)$

$\text{read}(B)$

$\text{write}(B)$

$(A)_1$

$T_1$

$\text{read}(A)$

$\text{write}(A)$

$\text{read}(B)$

$\text{write}(B)$

$R(A)$

$w(A)$

$r(B)$

$T_2$

$(A)_2$

$\text{read}(A)$

$\text{write}(A)$

$(A)_1$

$\text{read}(B)$

$\text{write}(B)$

$(A)_2$

$\text{read}(A)$

$\text{write}(A)$

$(A)_1$

$\text{read}(B)$

$\text{write}(B)$

$(A)_2$

$\text{r}(A)$

$\text{r}(B)$

$w(A)$

$w(B)$

$w(A)$

$R(A)$

$w(A)$

$r(B)$

$w(B)$

$r(A)$

$w(A)$

$r(B)$

$w(B)$

$\text{read}(B)$

$\text{write}(B)$

$(A)_1$

$R(A)$

$w(A)$

$r(B)$

$S_2$  is serial  
 $\rightarrow S_1$  is conflict

$S_1$  is serializable  
 $\rightarrow S_1$  is conflict

$S_1$  is equivalent to  $S_2$ .

serial schedule  
 $\text{read}(A)$   
 $\text{write}(A)$   
 $\text{read}(B)$   
 $\text{write}(B)$

serial schedule  
 $\text{read}(A)$   
 $\text{write}(A)$   
 $\text{read}(B)$   
 $\text{write}(B)$

No conflict

→ Concurrency control module generates schedules such that they don't produce any anomalies.

Check whether the following schedule is conflict serializable.

$S_1: r_2(A), w_2(A), r_1(A), w_1(A), r_2(B), w_2(B)$

$T_1$	$T_2$
	$r_2(A)$
	$w_2(A)$
$r_1(A)$	
$w_1(A)$	
	$r_2(B)$
	$w_2(B)$
$\Rightarrow r_2(B)$	$\Rightarrow r_1(A)$
$r_1(A)$	
$w_2(B)$	
$w_1(A)$	
$\Rightarrow r_2(B)$	
$r_1(A)$	
$w_2(B)$	
$r_1(B)$	
$w_2(B)$	
$r_1(A)$	
$w_1(A)$	

$\therefore S_2: r_2(A), w_2(A), r_2(B), w_2(B), r_1(A), w_1(A)$

$S_1$  and  $S_2$  are conflict equivalent and  $S_2$  is serial schedule

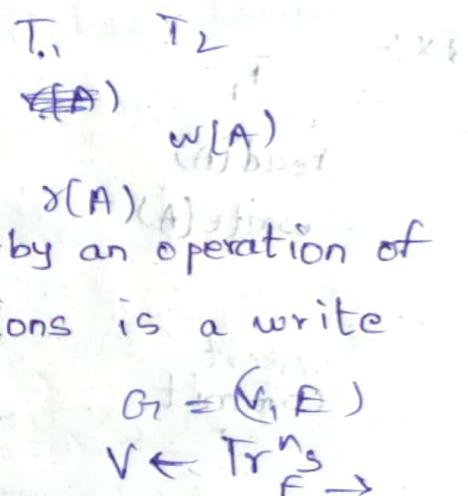
$\therefore S_1$  is conflict serializable

- There is a simple way to check whether a schedule is conflict serializable using a graph.
- The graph we construct is precedence graph.
- If the graph has a cycle then it is not conflict serializable.
- Nodes in the graph are transactions.
- Edges are determined by overlapping ops.

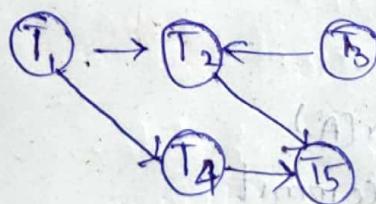
01/03/2022

Precedence graph:-

- If an operation  $T_i$  is followed by an operation of  $T_j$  and one of the two operations is a write then add edge from  $T_i$  to  $T_j$ .



S:  $w_1(A), r_2(A), w_1(B), w_3(C), r_{12}(C), r_4(B), w_2(D), w_4(E), r_5(D), w_5(E)$



→ As there is no cycle in precedence graph

→ The graph is conflict serializable.

→ To get the serial schedule, apply topographical sort for the directed graph.

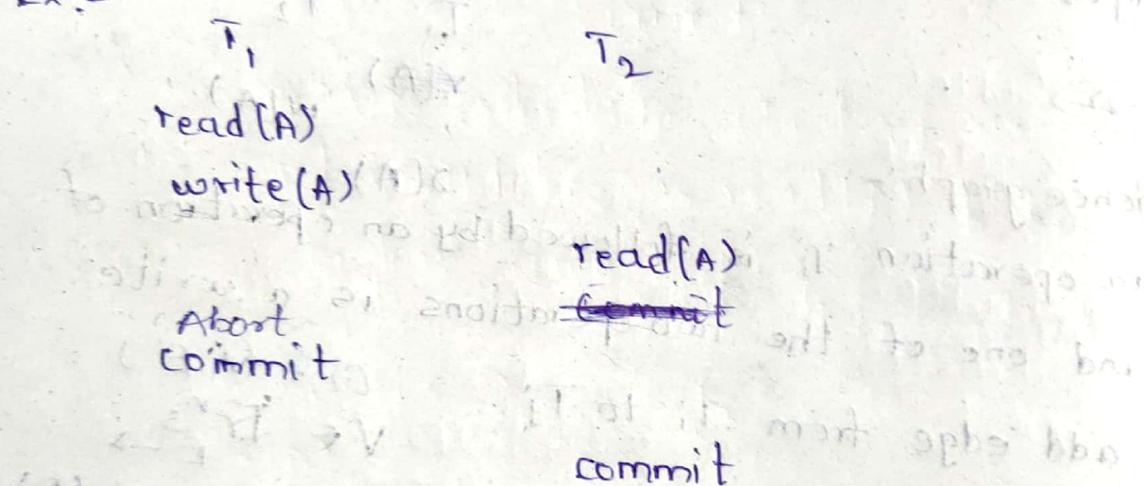
→ Suppose the two transactions  $T_i, T_j$

$T_i$  writes on item A : first

$T_j$  reads item A : second

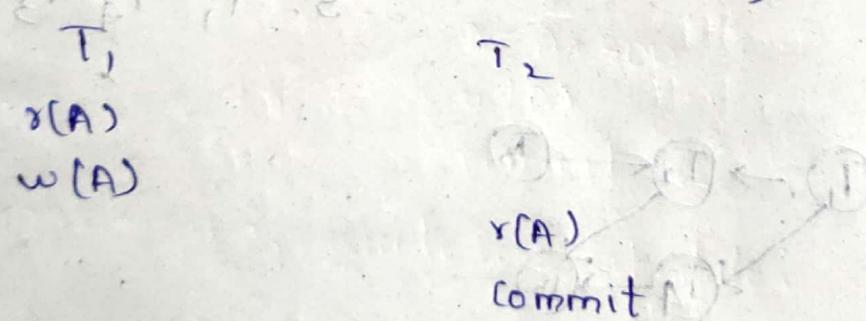
If commit of  $T_i$  proceeds before commit of  $T_j$ , then that schedule is called recoverable schedule.

Ex:-



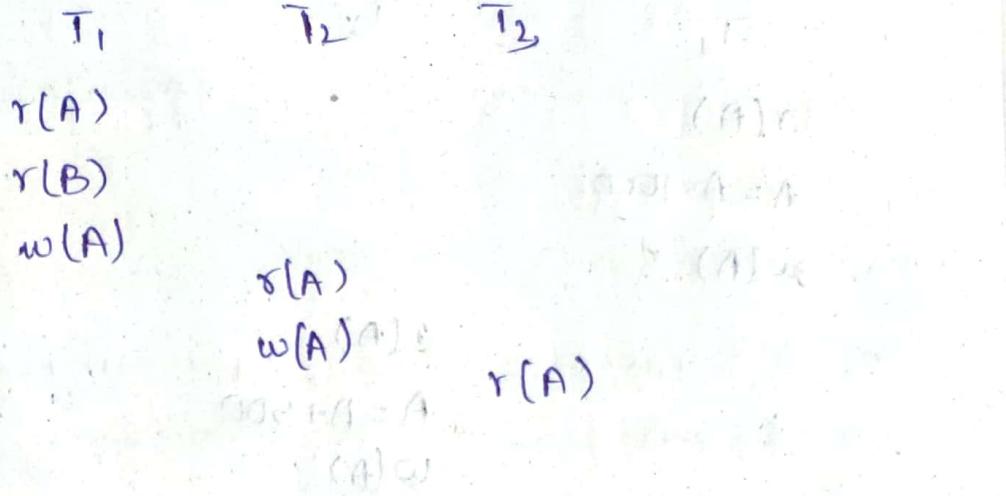
→ This is a recoverable schedule.

If the schedule is in this form,

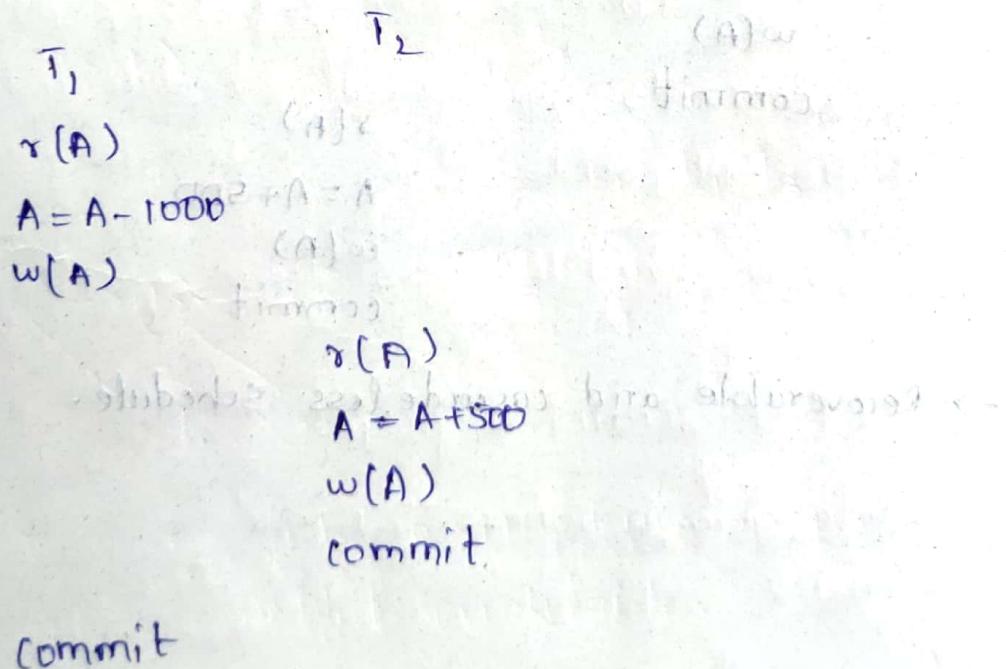


→ This is non-recoverable schedule.

If there is a role back in  $T_1$  after committing  $T_2$ ,  $T_2$  can't be rolled back.



- If T<sub>1</sub> aborts after  $r(A)$ , in T<sub>3</sub>, then all the three transactions roll back because of the recoverable schedule condition.
- This is called as cascade.
- If  $r(A)$  of T<sub>2</sub> is after commit of T<sub>1</sub> and if  $r(A)$  of T<sub>3</sub> is after commit of T<sub>2</sub> then schedule is cascade less schedule.



- This is non-recoverable schedule.

$T_1$  $r(A)$  $A = A - 1000$  $w(A)$  $r(A)$  $A = A + 500$  $w(A)$ 

commit

commit

- A recoverable schedule.
- It is not cascadeless as  $r(A)$  of  $T_2$  is happening before commit of  $T_1$ .

 $T_1$  $T_2$  $r(A)$  $A = A - 1000$  $w(A)$ 

commit

 $r(A)$  $A = A + 500$  $w(A)$ 

commit

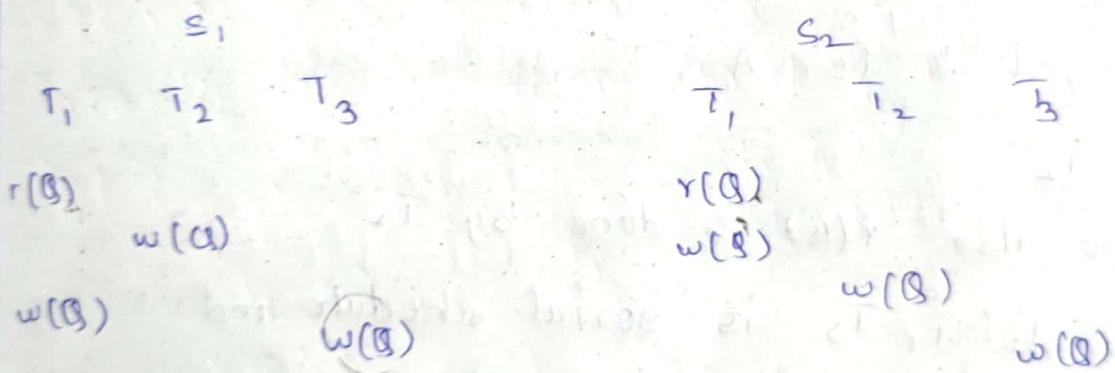
- Recoverable and cascadeless schedule.

(A)  
for  
from

10-03-2022

### View serializability:-

- If in schedule  $S$ ,  $T_i$  reads the initial value of  $Q$  then in  $S'$  also,  $T_i$  must read initial value of  $Q$ .
- If in schedule  $S'$ ,  $T_i$  executes  $\text{read}(Q)$  and that value was produced by  $T_j$  (if any) then in schedule  $S$  also  $T_i$  must read the value of  $Q$  that was produced by the same  $\text{write}(Q)$  operation of  $T_j$ .
- The transaction (if any) that performs the final  $\text{write}(Q)$  operation in  $S$  must also perform the final  $\text{write}(Q)$  operation in  $S'$ .



- In both schedules, the above conditions hold
- So,  $S_1$  and  $S_2$  are view equivalent
- As  $S_2$  is serial schedule,  $S_1, S_2$  are view serializable.
- View serializability is a restricted form of conflict serializability.
- If a write to an operand is performed directly without a read, then that write is called blind write.

→ If a schedule is view serializable, is not conflict serializable, then it should have at least one blind write.

Ex:-

S:  $R_2(B), R_2(A), R_1(A), R_3(A), W_1(B), W_2(B), W_3(B)$

B:  $T_2$

$T_3$

choose the schedule with starting as  $T_2$  and ending with  $T_3$

Our required serial schedule is  $T_2 T_1 T_3$

Now, check for A

A:  $T_2$

Here also  $R(A)$  is done by  $T_2$

So,  $T_2 T_1 T_3$  is serial schedule and satisfies all points.

So, given schedule S is view serializable.

Transaction control language.

- Delete from T<sub>1</sub> where T<sub>1</sub>.c = 'a'  
  Rollback

- Save point SP<sub>1</sub>  
- Delete from customer where Id = 1  
- Savepoint SP<sub>2</sub>  
  Delete from

Save point at

delete from customers where T1 = 3

Rollback to this

15-03-2023

### Concurrency Module Control

→ Concurrency is achieved by locking

→ Two types:  
(i) Shared lock (lock=S) → read

(ii) Exclusive lock (lock=X) → write

lock=S(A)      unlock=S(A)

read(A)      unlock=S(A)

unlock(A)      unlock=X(A)

→ If lock held is shared lock, then other transaction can read that element / can get that lock

→ If lock is exclusive lock, then other tr can't get that lock.

### Lock compatibility Matrix

	S	X
S	True	False
X	False	False

→ These locks don't ensure consistency. They provide isolation.

Ex:-

$T_1$ (A) reads	$\Rightarrow$	$T_2$
read(B)		lock-X(B)
B = B - 50		read(B)
write(B)		B = B - 50
r(A)		write(B)
		unlock(A)

$A = A + 50$

w(A)

$A = A + 50$

write(A)

unlock(A)

→ Converting a shared lock to exclusive lock is lock upgrading.

→ Converting exclusive lock to shared lock is lock downgrading.

$T_1$	$\Rightarrow$	$T_2$
r(A)		lock-S(A)
r(B)		r(A)
display(A+B)		unlock(A)
		lock-S(B)
		r(B)
		unlock(B)
		display(A+B)

→ Let the transactions interleave this way;

$T_1$

lock-X(B)  
read(B)  
 $B = B - 50$   
write(B)  
unlock(B)

$T_2$

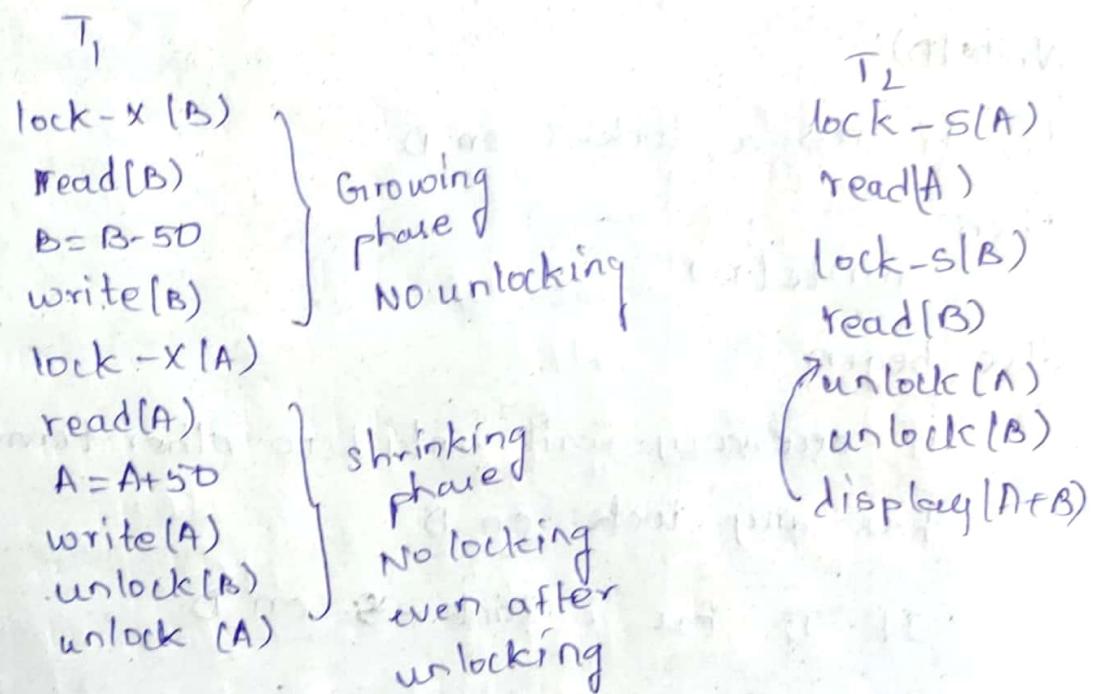
lock-S(A)  
read(A)  
unlock(A)  
lock-S(A)  
read(B)  
unlock(B)  
display(A+B)

lock-X(A)  
read(A)  
 $A = A + 50$   
write(A)  
unlock(A)

- Here, when  $A = 100$ ,  $B = 200$ , the exp res is 300, but, the actual res is 250  
→ So, lock solemnly doesn't ensure consistency.

### Two-phase locking protocol (2PL)

- It says that all the locks should be attained and then only unlocking should start.



17-03-2022

Read(D):

If  $T_i$  has a lock on D  
then

    read(D)

else

begin

    if necessary, wait until no other transaction

        has a lock-X on D

        grant  $T_i$  a lock-S on D

    read(D)

end

Write(D):

If  $T_i$  has a lock-X on D

then

    write(D)

else begin

    if necessary wait until no other transaction

        has any lock on D

    If  $T_i$  has a lock-S on D

        then

            upgrade lock on D to lock-X

        else grant a lock-X on D

    write(D)

end

After commit() / Abort unlock

→ The point of time at which last lock was obtained by the transaction is called lock-point of transaction.

Theorem:-

Two phase locking produces only conflict serializable schedule.

Proof:-

Assume  $T_0, T_1, T_2, \dots, T_{n-1}$  be transactions in the schedule.

Let us assume there is a cycle in precedence graph involving all tr<sup>n</sup>s.

If there is an  $T_i \rightarrow T_j$  in graph, then

$$\alpha_i < \alpha_j$$

so, if  $T_0 \rightarrow T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_{n-1} \rightarrow T_0$   
 $\Rightarrow \alpha_0 < \alpha_1 < \alpha_2 < \dots < \alpha_{n-1} < \alpha_0$   
 $\Rightarrow \alpha_0 < \alpha_0$

which is contradiction.

So, there will not be any conflict.

Hence two phase locking gives only conflict serializable schedule.

Deadlock:-

$T_1$   
lock-X(B)

r(B)

$$B = B + 50$$

w(B)

$T_2$

lock-S(A)

r(A)

lock-S(B)

lock-X(A)

→  $T_1$  waits for lock on A

$T_2$  waits for lock on B

→ So, a deadlock

→ We should one of the tr's to deal with deadlock.

$T_1$

$T_2$

$T_3$

t-X(A)

r(A)

l-S(B)

r(B)

w(A)

unlock(A)

lock-X(A)

r(A)

w(A)

u-I(A)

l-S(A)

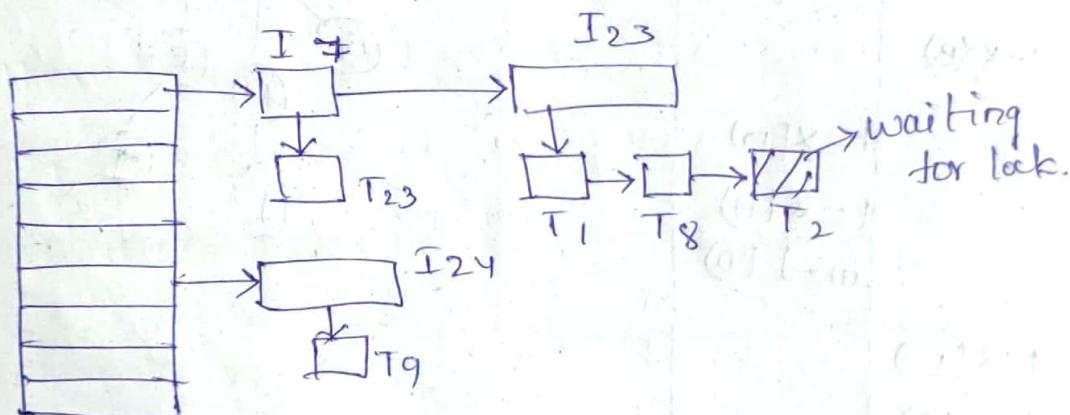
r(A)

→ If there is a deadlock in above schedule then if we decide to rollback  $T_1$ , then  $T_2$  &  $T_3$  should also be rolled back due to cascading.

→ A transaction, if it is waiting for the lock for long time, it is called starvation.

- A transaction if it holds all its exclusive locks till it commits/aborts, then cascading can be avoided. "strict 2 phase" rollback.
- A transaction if it holds all the locks till it commits or aborts.
- This is rigorous 2 phase locking
- The d.s used for implementing locks is lock-table.

Lock Table:-



22-03-2022

Graph Based Protocol

- 1) Only exclusive locks are allowed.
- 2) The first lock by  $T_i$  may be on any data item subsequently a data  $Q$  can be locked by  $T_i$  only if the parent of  $Q$  is currently locked by  $T_i$ .
- 3) Data items may be unlocked at any time.
- 4) A data item that has been locked and unlocked by  $T_i$  cannot be subsequently be locked by  $T_i$ .

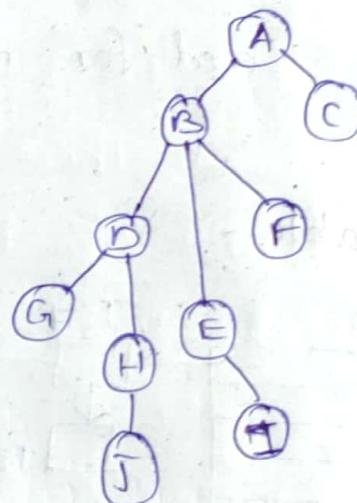
Ex:-

T<sub>0</sub>: lock-x(B), lock-x(E), lock-x(D), unlock(B), unlock(E), lock-x(H), unlock(D), unlock(G)

T<sub>1</sub>: l-x(D), l-x(H), unlock(D), unlock(H)

T<sub>2</sub>: l-x(B), l-x(E), unlock(E), unlock(B)

T<sub>3</sub>: l-x(D), l-x(H), unlock(D), unlock(H)



T <sub>0</sub>	T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
l-x(B)			
	l-x(D)		
	l-x(H)		
	un-l(D)		
l-x(E)			
l-x(D)			
un-l(B)			
un-l(E)			
	lock-x(B)		
	l-x(E)		
	unlock(H)		
l-x(G)			
un-l(D)			
		l-x(D), l-x(H)	
		unlock(D)	
		unlock(H)	
		<del>unlock(I)</del>	
		u-l(E)	
		u-l(B)	
u-l(G)			

→ It ensures that the schedule is free from deadlock but doesn't guarantee recoverability and cascade freedom.

Commit Dependency:-

→  $\emptyset$  is a data item with uncommitted write.

→ Record which transaction performed last write to  $\emptyset$  ( $T_j$ )

→  $T_i$  reads  $\emptyset$

→  $T_i$  is not allowed to commit until  $T_j$  commits.

Deadlock Handling:-

1) Deadlock Prevention

2) Deadlock Detection

Deadlock Prevention:-

→ acquire all locks before starting the transaction.

→ But concurrency suffers.

→ Impose partial ordering for all data items before it begins execution.

→ Require that a transaction can lock data items only in order specified by the partial order.

Timeout schemes:-

→ If any lock is not obtained, then wait for loops and retry the lock for acquiring.

(23-03-2022)

## Deadlock Prevention:-

- wait - die scheme
- older transactions wait for younger transactions to release data item.

Younger transactions never wait for older ones, they are rolled back.

Old - wait

Young - die.

- wound - wait scheme
- older tr's forces roll back of younger tr's instead of waiting for it.

Younger transactions wait for older ones.

wound - older tr's (kills younger tr's)

wait - younger tr's (wait for older to release)

→ The above schemes may result in starvation.

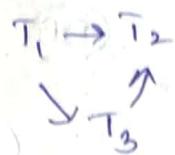
→ To avoid starvation, make sure that oldest tr is not rolled back.

## Deadlock Detection:-

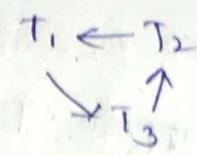
→ If  $T_i$  waits for  $T_j$  to release a lock, then draw an edge from  $T_i$  to  $T_j$

→ If there is a cycle in the graph, then deadlock has occurred.

$$T_i \rightarrow T_j$$



No deadlock  
No cycle



Deadlock  
(has cycles)

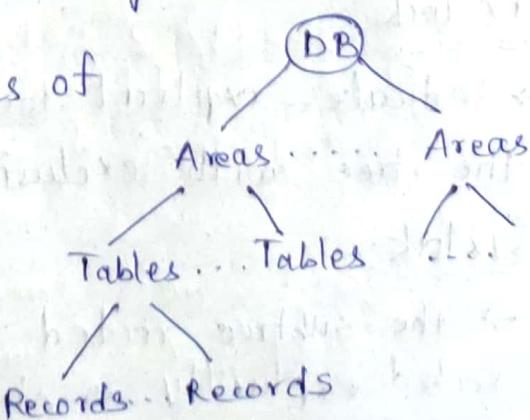
### Deadlock Recovery:-

- kill one of the transaction causing the deadlock
- This is the only way to recover from deadlock.
- Two types of rollback:-
- 1) Total Rollback:- Blindly rollback entire trn.
- 2) Partial Rollback:- Rollback till certain instrn.

### Multiple Granularity:-

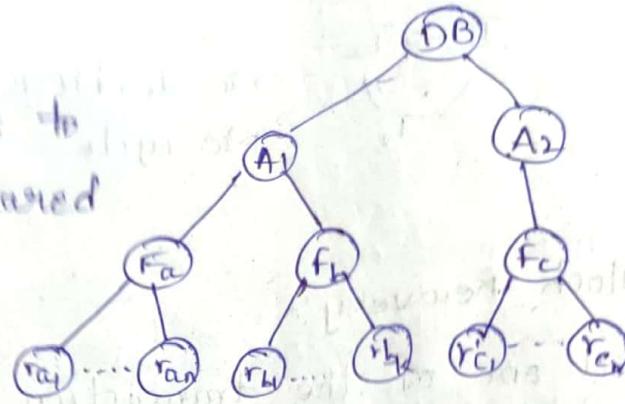
- If we have to update all records in the table to a single value, then instead of obtaining a lock for each record, acquire a single lock for the table.
- If we consider only few records, obtain a lock for each record.
- Obtaining locks at various levels of db is called multiple granularity.

- These are various levels of the database.



Ex:-

Suppose  $T_1$  wants to read  $r_{a_2}$  and shared lock is supplied.



At that time,  $T_2$  wants to write to db, but ex-lock is not acquired due to shared lock on  $r_{a_2}$ .  
Also, it should traverse across all nodes to  
→ This problem is due to multiple check if x-lock can be granted.

→ There are locks called intention locks

1) IS lock

(Intention shared lock)

2) IX lock

(Intention exclusive lock)

3) SIX lock (Intention shared - exclusive lock)

IS lock:-

→ Indicates explicit locking at a lower level of the tree but with shared locks.

IX lock

→ Indicates explicit locking at a lower level of the tree with exclusive locks / shared locks.

SIX lock:-

→ The subtree rooted by the present node is locked explicitly in shared mode and explicit

locking is done at a lower level with X-locks.

(24-03-2022)

Intention lock compatibility matrix:-

	IS	IX	S	SIX	X
IS	T	T	T	T	F
IX	T	T	F	F	F
S	T	F	T	F	F
SIX	T	F	F	F	F
X	F	F	F	F	F

Rules for  $T_i$  to lock a node  $Q$ :

- 1) The lock compatibility matrix should be considered.
- 2) The root of tree must be locked first and may be locked in any mode.
- 3) A node  $Q$  can be locked by  $T_i$  in S (or) IS mode, the parent of  $Q$  is currently locked by  $T_i$  in either IX (or) SIX mode.
- 4) A node  $Q$  can be locked by  $T_i$  in X, SIX, IX mode only if the parent of  $Q$  is currently locked by  $T_i$  in either IX (or) SIX mode.
- 5) No unlocking.
- 6) All the locks on its children must be released before releasing locks on particular node.

Index locking

select count(\*) from ins where dept = Physics  
insert into ins values  
(a, b, Physics, c) → req X

↳ req shared

→ But there is no conflict b/w them.

→ This is Phantom phenomenon.

→ This problem can be solved by acquiring a lock on data item dept\_name!

→ Or use index locking.

### Time stamp based protocol

→ Two timestamps for any item

- 1) Read timestamp → largest T.S of T; which made read
- 2) write timestamp  
↳ largest T.S among tr's that made write (Q)

(30-03-2022)

→ Time stamp is a counter.

→ Every transaction is assigned a time stamp when it arrives in the system.

→ If  $T_i$  arrives before  $T_j$ , then

$$TS(T_i) < TS(T_j)$$

→ For a data item, it is provided with two time stamps.

write - w - timestamp (Q) -  $w \cdot TS(Q)$

read - r - timestamp (Q) -  $r \cdot TS(Q)$

where  $r\text{-TS}(\varnothing)$  = Time stamp of latest transaction which successfully read  $\varnothing$ .  
 $w\text{-TS}(\varnothing)$  = Time stamp of latest transaction which successfully wrote  $\varnothing$ .

→ Timestamp order = serializability order.

Read( $\varnothing$ ):-

→  $T_i$  wants to read  $\varnothing$ .

→ Time stamp of  $T_i$  is  $TS(T_i)$ .

if  $TS(T_i) \leq w\text{-Timestamp}(\varnothing)$

then  $T_i$  is rolled back

→ As the value of  $\varnothing$  is overwritten by a transaction that has a ts greater than  $T_i$ .

if  $TS(T_i) > w\text{-Timestamp}(\varnothing)$

then  $\text{read}(\varnothing)$  is successfully done by  $T_i$ .

$$R\text{-TS}(\varnothing) = \max(R\text{-TS}(\varnothing), TS(T_i))$$

Write( $\varnothing$ ):-

→  $T_i$  wants to write  $\varnothing$ .

If  $TS(T_i) < R\text{-Timestamp}(\varnothing)$

then  $T_i$  is rolled back

If  $TS(T_i) < w\text{-Timestamp}(\varnothing)$

then  $T_i$  is rolled back.

otherwise  $T_i$  can write  $\varnothing$ .

$$w\text{-Timestamp}(\varnothing) = TS(T_i)$$

Ex-1 (contd) To generate initial value problem

$T_1$	$\theta \leftarrow T_2$	Initial cond:-
$r(T)$	$w(T)$	$TS(T_1) = 1$
		$TS(T_2) = 2$
$w(T)$		$w - TS(\theta) = 0$
		$r - TS(\theta) = 0$

1) Read( $\theta$ )

$$TS(T_1) = 1$$

$$w - TS(\theta) = 0$$

$$TS(T_1) > w - TS(\theta)$$

$$T_1 \xrightarrow{\text{reads}} \emptyset$$

$$R - TS(\theta) = 1$$

2) write( $\theta$ )

$$TS(T_2) = 2$$

$$R - TS(\theta) = 1$$

$$w - TS(\theta) \approx 0$$

$$TS(T_2) > r - TS(\theta)$$

$$TS(T_2) > w - TS(\theta), T_2 \text{ writes } \theta, w - TS(\theta) = 2$$

3) write( $\theta$ )

$$TS(T_1) = 1$$

$$R - TS(\theta) = 1$$

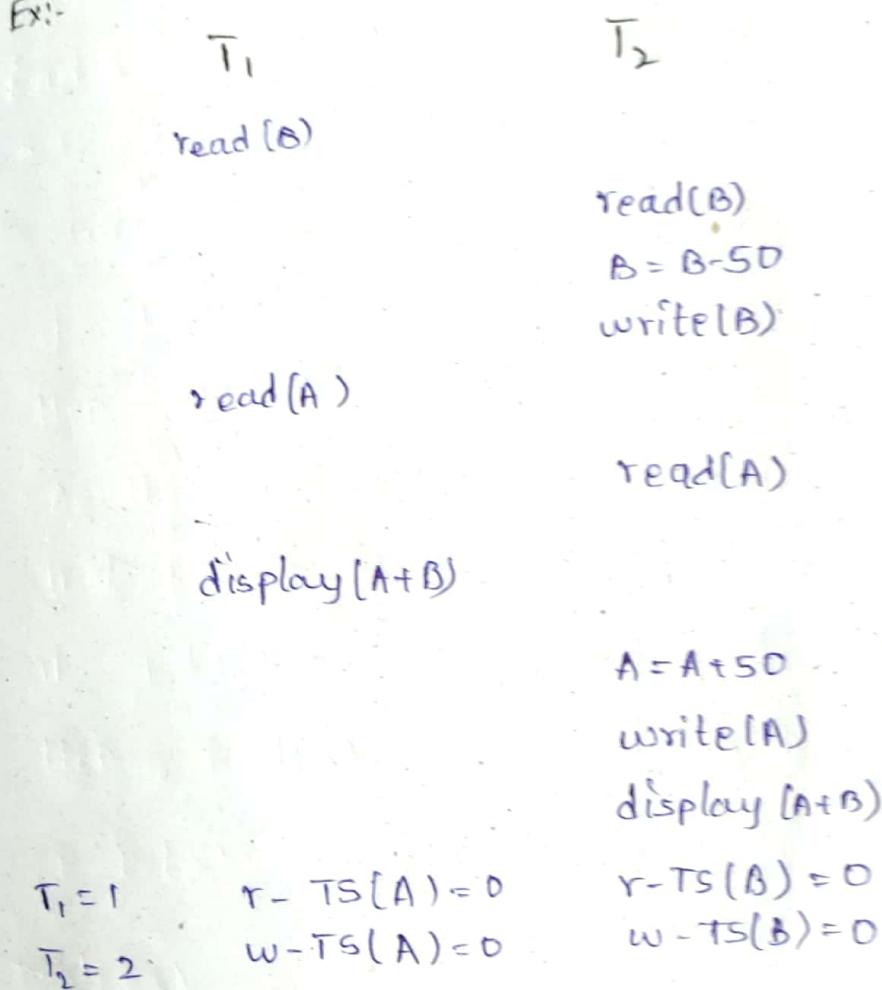
$$w - TS(\theta) = 2$$

$$TS(T_1) < w - TS(\theta)$$

$T_1$  roll back

- If there is a obsolete write, then ignore that write and proceed with other instns.
- This rule is called Thomas write rule.
- This produces view serializable schedules.

Ex:-



1)  $r(B)$

$$TS(T_1) = 1, w - TS(B) = 0$$

$$TS(T_1) > w - TS(B)$$

$T_1$  reads  $B$ ,  $r - TS(B) = 1$

2)  $r(B)$

$$TS(T_2) = 2, w - TS(B) < 0$$

$$TS(T_2) > w - TS(B)$$

$T_2$  reads  $B$ ,  $r - TS(B) = 2$

3)  $w(B)$ ,  $TS(T_2) = 2, w - TS(B)$